

Feb 4  
Consistent Global States  
(Part 2)

Vijay Chidambaram

# Consistent observations

- A consistent observation is one that corresponds to a consistent run.
- It is the possibility of messages being reordered by channels that leads to undesirable observations
- We can restore order to messages between pairs of processes by defining a delivery rule for deciding when received messages are to be presented to the application process.



# FIFO Delivery

- FIFO Delivery:  $\text{sendi}(m) \rightarrow \text{sendi}(m') \Rightarrow \text{deliveri}(m) \rightarrow \text{deliveri}(m')$
- Clock Condition:  $e \rightarrow e' \Rightarrow \text{Timestamp}(e) < \text{timestamp}(e')$
- Delivery Rule DR1: At time  $t$ , deliver all received messages with timestamps up to  $t - \text{delta}$  in increasing timestamp order

# Gap Detection

- Given two events  $e$  and  $e'$  along with their clock values  $LC(e)$  and  $LC(e')$  where  $LC(e) < LC(e')$ , determine whether some other event  $e''$  exists such that  $LC(e) < LC(e'') < LC(e')$
- A message  $m$  received by process  $p$  is called stable if no future messages with timestamps smaller than  $TS(m)$  can be received by  $p$ .
- DR2: Deliver all received messages that are stable at  $p_0$  in increasing timestamp order.



# Causal Delivery

- FIFO delivery is per-process
- Causal Delivery extends it across processes
- Causal Delivery:  $\text{send}_i(m) \rightarrow \text{send}_j(m') \Rightarrow \text{deliver}_i(m) \rightarrow \text{deliver}_j(m')$ 
  - Note  $i$  and  $j$  can be different
- FIFO delivery between all pairs of processes not enough for causal delivery
- $A \rightarrow B, A \rightarrow C; B \rightarrow C$ 
  - $\text{send}_{AB}(M1) \rightarrow \text{send}_{AC}(M2)$
  - B gets M1, sends it to C
  - C gets M2 before M1.
  - FIFO delivery is ensured, causal delivery is not.

# Consistent Observations

- If the observer process uses a delivery rule that satisfies Causal Delivery, then all its observations will be consistent
- How do we build this in a practical distributed system?



# Building the Observer

- Assume Observer has two parts: the Logic Controller and the Network Controller
- The Network Controller gets events sent by other nodes, decides in what order to show them to Logic Controller
  - Delivery rules are implemented here
- Logic Controller takes actions based on what it sees ("declare deadlock")

# Building the Observer

- Network Controller gets two events E1 and E2
- Timestamp = TS
- $TS(E1) < TS(E2)$
- This doesn't mean  $E1 < E2$
- Only this is true:  $E1 < E2 \Rightarrow TS(E1) < TS(E2)$
- So we know  $E2 < E1$  is false, E1 and E2 could still be concurrent



# Strong Clock Condition

- Need stronger condition to implement Network Controller
- Strong Clock Condition:
  - $E1 < E2 \Rightarrow TS(E1) < TS(E2)$
  - $TS(E1) < TS(E2) \Rightarrow E1 < E2$

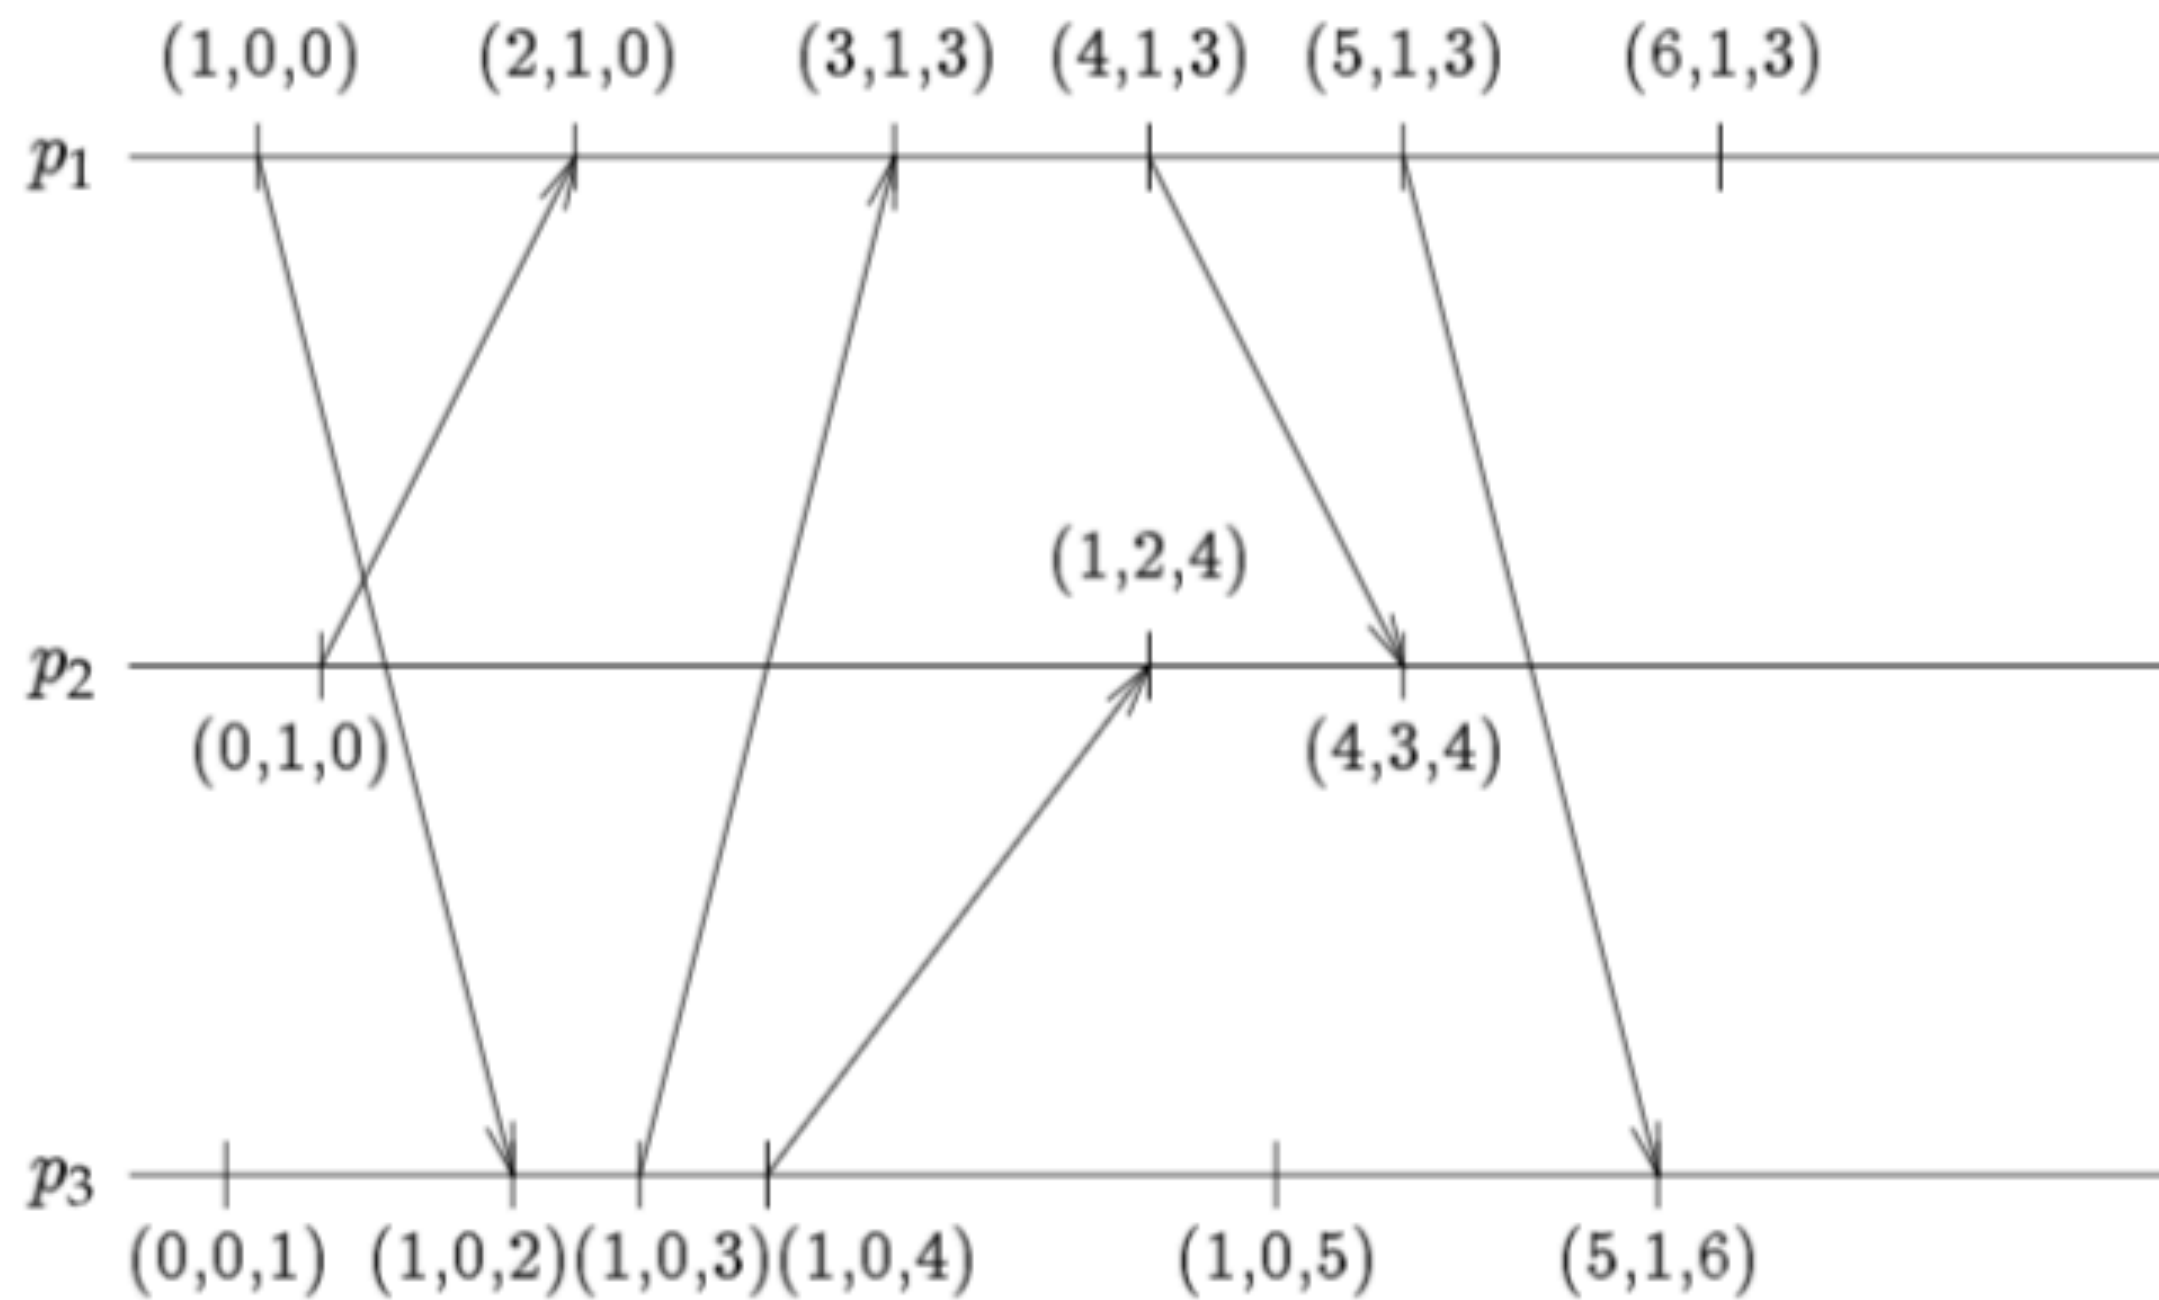
# Implementing the Strong Clock Condition

- Brute force approach:
  - At each node, keep the Causal History  $C$
  - Causal History  $C(e)$  is the set of all events  $e'$  such that  $e' < e$
  - All previous local events at node are part of  $C(e)$
  - When  $A$  sends a message to  $B$ ,  $C(\text{receipt of message } M) = C(\text{previous local event at } B) \cup C(\text{sending event at } A)$
  - $E1 < E2$  if  $C(E1)$  is a subset of  $C(E2)$
- Problem: the set  $C$  grows too large too quickly, impractical to maintain



# Vector Clocks

- Each node maintains a vector  $V[n]$  where there are  $n$  nodes
- $V = 0$  on initialization for all nodes
- For local event  $E_i$  at node  $i$ , update  $V[i] = V[i] + 1$
- On receipt( $M$ ),
  - $V = \max(V, \text{vector-TS}(M))$  for each element of  $V$
  - $V[i] += 1$
- $V[j]$  = number of events of  $j$  that casually precede this event (when this process is not  $j$ )
- $V[i]$  = number of local events when this process is  $i$





# Using Vector Clocks

- $V1 < V2$  if no element of  $V2$  is bigger than its corresponding element in  $V1$
- Strong Clock Condition:  $E1 < E2 \iff V(E1) < V(E2)$
- $V1$  and  $V2$  are concurrent if  $V1[i] > V2[i]$  for some  $i$  and  $V2[j] > V1[j]$  for some  $j$

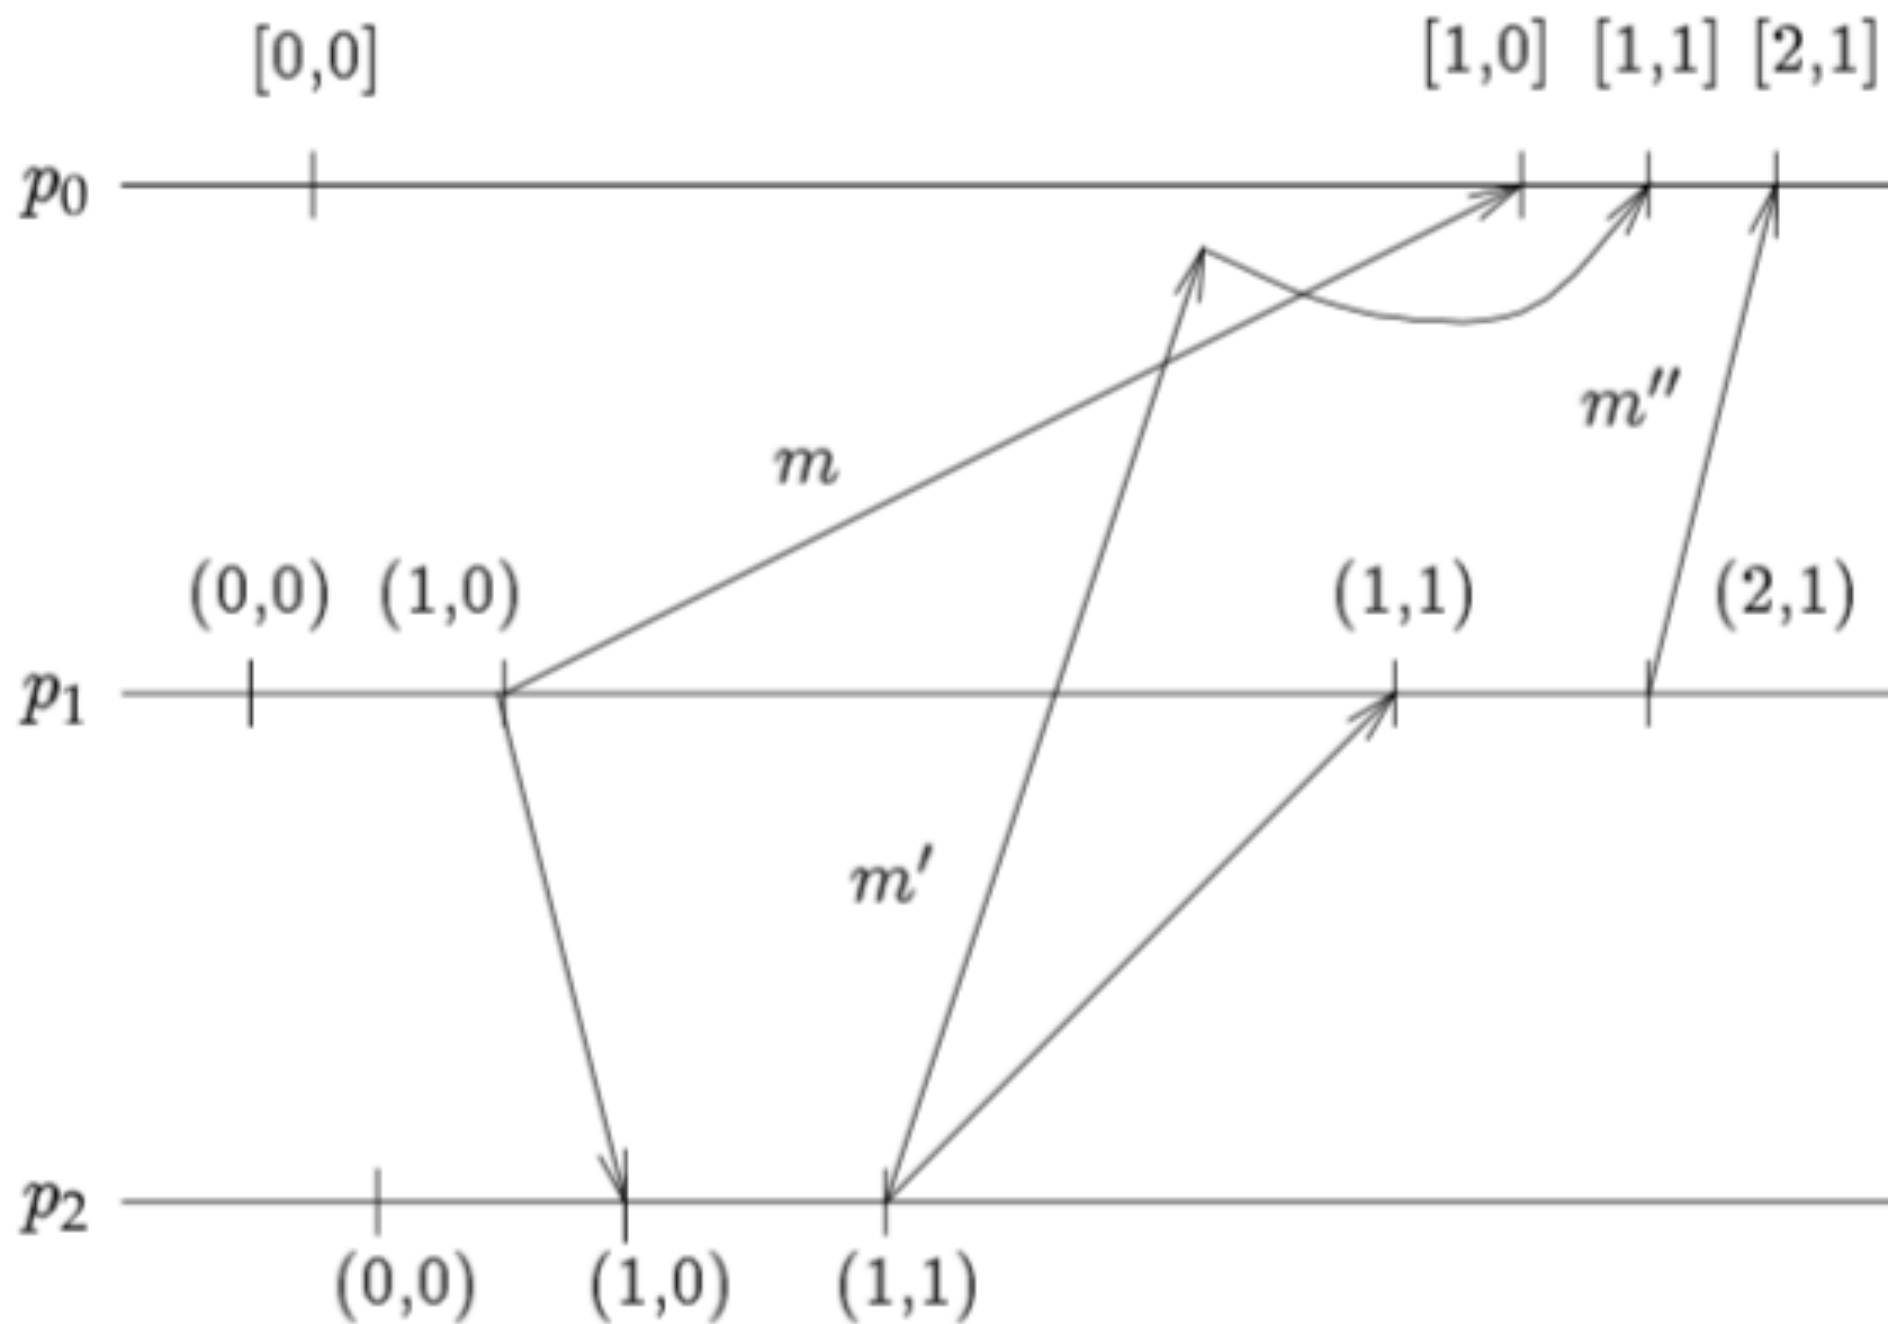
# Using Vector Clocks in the Network Observer

- All events sent to observer have vector timestamps
- Network Observer obtains set of events  $M$ , then decides on order to deliver them
- A message can be delivered as soon as Network Observer determines there are no events that casually precede this message
- Consider message  $M$  from  $J$
- $M'$  is last message delivered from  $K$  ( $K \neq J$ )
- To deliver  $M$ , NO has to determine that:
  - There is no earlier message from  $J$  that is undelivered
    - if  $V(M)[J] - 1$  message have already been delivered, there cannot be an earlier message
  - There is no undelivered message  $M''$  such that  $\text{send}_K(M') < \text{send}(M'') < \text{send}_J(M)$ 
    - $V(M')[K] \geq V(M)[K]$  for all  $K \Rightarrow V(M) < V(M')$



# Using Vector Clocks in the Network Observer

- Network Observer maintains array  $D$ , initialized to 0
- Deliver message  $M$  with time stamp  $V$  from  $J$  as soon as:
  - $D[J] = V[J] - 1$
  - $D[K] \geq V[K]$ , for all  $K \neq J$
- When Network Observer delivers  $M$ ,  $D$  is updated by setting  $D[j] = V[j]$



**Figure 8. Causal Delivery Using Vector Clocks**