

Jan 30

Consistent Global States of  
Distributed Systems: Fundamental  
Concepts and Mechanisms

Vijay Chidambaram

# Global State of a distributed system

- Union of local state of nodes
- No way to get instantaneous snapshot of all nodes
- Only way to get local state of a node is by sending it a message
- How to find a meaningful global state of the system?



# Why is this hard?

- Messages could be dropped or delayed
- Computed Global state may be:
  - Obsolete: state changed since we have checked
  - Incomplete: state of some nodes may be missing
  - Inconsistent: imagine a token is sent from A to B. Computed Global state may show token in both A and B

# Computing Global State

- What if we simply used a lot of messages to nodes to compute global state?
- Would help with ensuring global state is complete or current (not obsolete), but will not help with ensuring global state is consistent
- Consistency cannot be achieved by throwing more messages at the problem



# Global Predicate Evaluation (GPE)

- We construct a predicate based on the global state
- We want to evaluate whether this global predicate is true or false
- Examples: the system is deadlocked, a majority of nodes are alive and responsive

# Modeling the system

- N sequential processes  $p_1.. p_n$
- Each pair of processes has a channel
- Channels are reliable but may deliver messages out of order
- Why do we model system as async?
  - Physical delays are bounded
  - Software creates unbounded delays



# Distributed Computation

- Activity distributed among  $N$  processes
- Each process sees three kinds of events:
  - Events local to that process
  - Send message to process  $p_i$
  - Receive message from process  $p_j$
- All events are recorded in the local history of the process
- Global history is union of histories of all participating processes

# Global History

- Global history does not order all events
- An event  $A$  is only ordered with respect to event  $B$  if  $A$  happening affects  $B$  in some way
- Events are ordered using “happens-before” relationships



# Lamport Clocks

- Notation:  $e(i, k)$  =  $k$ th event in process  $i$
- $e(i, j) < e(i, k)$  if  $j < k$
- if  $e(i, j) = \text{send}(m)$ , and  $e(k, l) = \text{receive}(m)$ 
  - then  $e(i, j) < e(k, l)$
- if  $e(i, j) < e(k, l)$  and  $e(k, l) < e(m, n)$ 
  - then  $e(i, j) < e(m, n)$
- All other events are considered concurrent
  - consider  $e(i, j)$  and  $e(k, l)$  concurrent
  - $e(i, j) < e(k, l)$  is false
  - $e(k, l) < e(i, j)$  is also false

# Distributed Computation

- Formally, a distributed computation is a partially ordered set (poset) defined by the pair  $(H, \rightarrow)$
- Not all events are ordered
- Events inside each process are totally ordered
- Events across processes are partially ordered



# Cuts

- A cut of a distributed computation is a subset  $C$  of its global history  $H$  and contains an initial prefix of each of the local histories
- A cut can be defined by tuple  $(c_1, c_2, \dots, c_N)$ 
  - Process  $p_i$ 's last event in the cut is  $c_i$
  - $P_1$ 's last event in the cut is  $c_1$
- Frontier of the cut: the set of events  $e(i, c_i)$  for  $i=1..n$  (the last events included in the cut for each process)




Figure 2. Cuts of a Distributed Computation



# Runs

- A run of a distributed computation is total ordering  $R$  that includes all of the events in the global history and that is consistent with each local history.
- For process  $P_i$ , the events of  $P_i$  occur in the same order in  $R$  as in the history of  $P_i$
- There are many possible runs for a single distributed computation with history  $H$

# Inconsistent Cuts

- Not all cuts are consistent
- If a cut includes receipt of a message but not the sending of the message, it is inconsistent
- More precisely, if  $e(i, j) < e(k, l)$  and  $e(k, l)$  is in the cut, then  $e(i, j)$  must also be in the cut
- Global properties must be checked using consistent cuts (which lead to consistent global states)
- Using inconsistent cuts may lead to determination of "ghost deadlock"



# Reachable States

- A run  $R$  is said to be consistent if for all events,  $e_1 < e_2$  implies that  $e_1$  appears before  $e_2$  in  $R$
- Each (consistent) global state  $S_i$  of the run is obtained from the previous state  $S_{i-1}$  by some process executing the single event  $e_i$
- $S_{i-1}$  leads to  $S_i$
- $S_j$  is reachable from  $S_i$ , if there is a series of consistent states from  $S_i$  to  $S_j$  such as  $S_i \rightarrow S_k \rightarrow S_j$

# Lattice

- The set of all consistent global states of a computation along with the leads-to relation defines a lattice.
- The lattice consists of  $n$  orthogonal axes, with one axis for each process
- Each path down the lattice is one run of the distributed system






Figure 3. A Distributed Computation and the Lattice of its Global States

# Observing a distributed system

- Idea 1: Active Observer
  - One monitor process sends messages to all processes
  - Constructs global state based on responses
  - Can lead to inconsistent cut



# Observing a distributed system

- Idea 2: Passive Observer
  - One monitor process gets copy of all messages sent by processes
  - Different monitors observe different cuts (and hence different global states)
  - Consistent observation leads to a consistent run

# Observing a distributed system

- We ensure messages from the same process are delivered in order (FIFO delivery)
  - Implemented using per-process sequence numbers
- Delivery Rule:
  - if  $e1 < e2$ , then  $ts(e1) < ts(e2)$ .
  - Deliver messages in timestamp order