

Feb 27

# Impossibility Result and Consensus

Vijay Chidambaram

# Consensus

- A set of asynchronous processes have to agree on a value
- In the more general case, each process proposes a value, and one of the values is agreed upon
- Simpler case is when the values are 0 or 1
- Some processes may fail
- Desired properties: termination, agreement, validity
- In FLP proof, we only need some process to decide (not all) – weak termination



# Impossibility Result

- No completely asynchronous consensus protocol can tolerate even a single unannounced process death.
- Assumptions:
  - no byzantine failures (only fail-stop)
  - reliable communication channels (exactly once reliably delivery, not FIFO, may be delayed and re-ordered)
  - processes do not have access to a synchronized clock

# Modeling the system

- Each process is a finite state automaton
- In one atomic step:
  - each process receives one message
  - does local computation in responses to message
  - sends out finite number of messages (atomic broadcast)



# Window of vulnerability

- Period of time in commit or consensus protocol when the processes wait indefinitely for a failed process
- Impossibility result indicates every fully async commit or consensus protocol has this window of vulnerability

# Modeling the system (more detail)

- A consensus protocol  $P$  is an asynchronous system of  $N$  processes ( $N \geq 2$ )
- Each process  $p$  has a one-bit input register  $x_p$ , an output register  $y_p$  with values in  $\{b, 0, 1\}$ , and an unbounded amount of internal storage
- Decision states: when output register has 0 or 1
- Output cannot change once decided



# Modeling the system (more detail)

- Message format: (destination, message-value)
- send() and receive() similar to Project 1
- Configuration: internal state of each process + message buffers
- A consensus protocol is partially correct if it satisfies two conditions:
  - (1) No accessible configuration has more than one decision value.
  - (2) For each  $v \in (0, 1)$ , some accessible configuration has decision value  $v$ .

# Correctness Condition

- A consensus protocol  $P$  is totally correct in spite of one fault if it is partially correct, and every admissible run is a deciding run.
- Our main theorem shows that every partially correct protocol for the consensus problem has some admissible run that is not a deciding run.



# Overview of proof

- The basic idea is to show circumstances under which the protocol remains forever indecisive.
- First, we argue that there is some initial configuration in which the decision is not already predetermined.
- Second, we construct an admissible run that avoids ever taking a step that would commit the system to a particular decision.

# Step 1: There is a bivalent initial config

- Consider all initial configs: some lead to 0, some lead to 1
- Order all initial configs such that two configs which differ only in one value are next value
- Somewhere in this initial config, there must be a pair where one config is 0, another config is 1
  - These two configs differ only by the value of one process (lets call that  $p$ )
- We allow one process to fail, so 0-config should decide 0 even if  $p$  fails
  - But if it does not depend on  $p$ , 1-config (which only differs by  $p$ ) should also decide 0.
  - Thus, 1-config is actually bivalent: it can decide 0 or 1



# Step 2

- If you start from a bivalent config and apply message  $e$ , the resulting set of configs contains a bivalent config
- In other words, if you delay a message  $e$  long enough, you can move from a bivalent config to another bivalent config

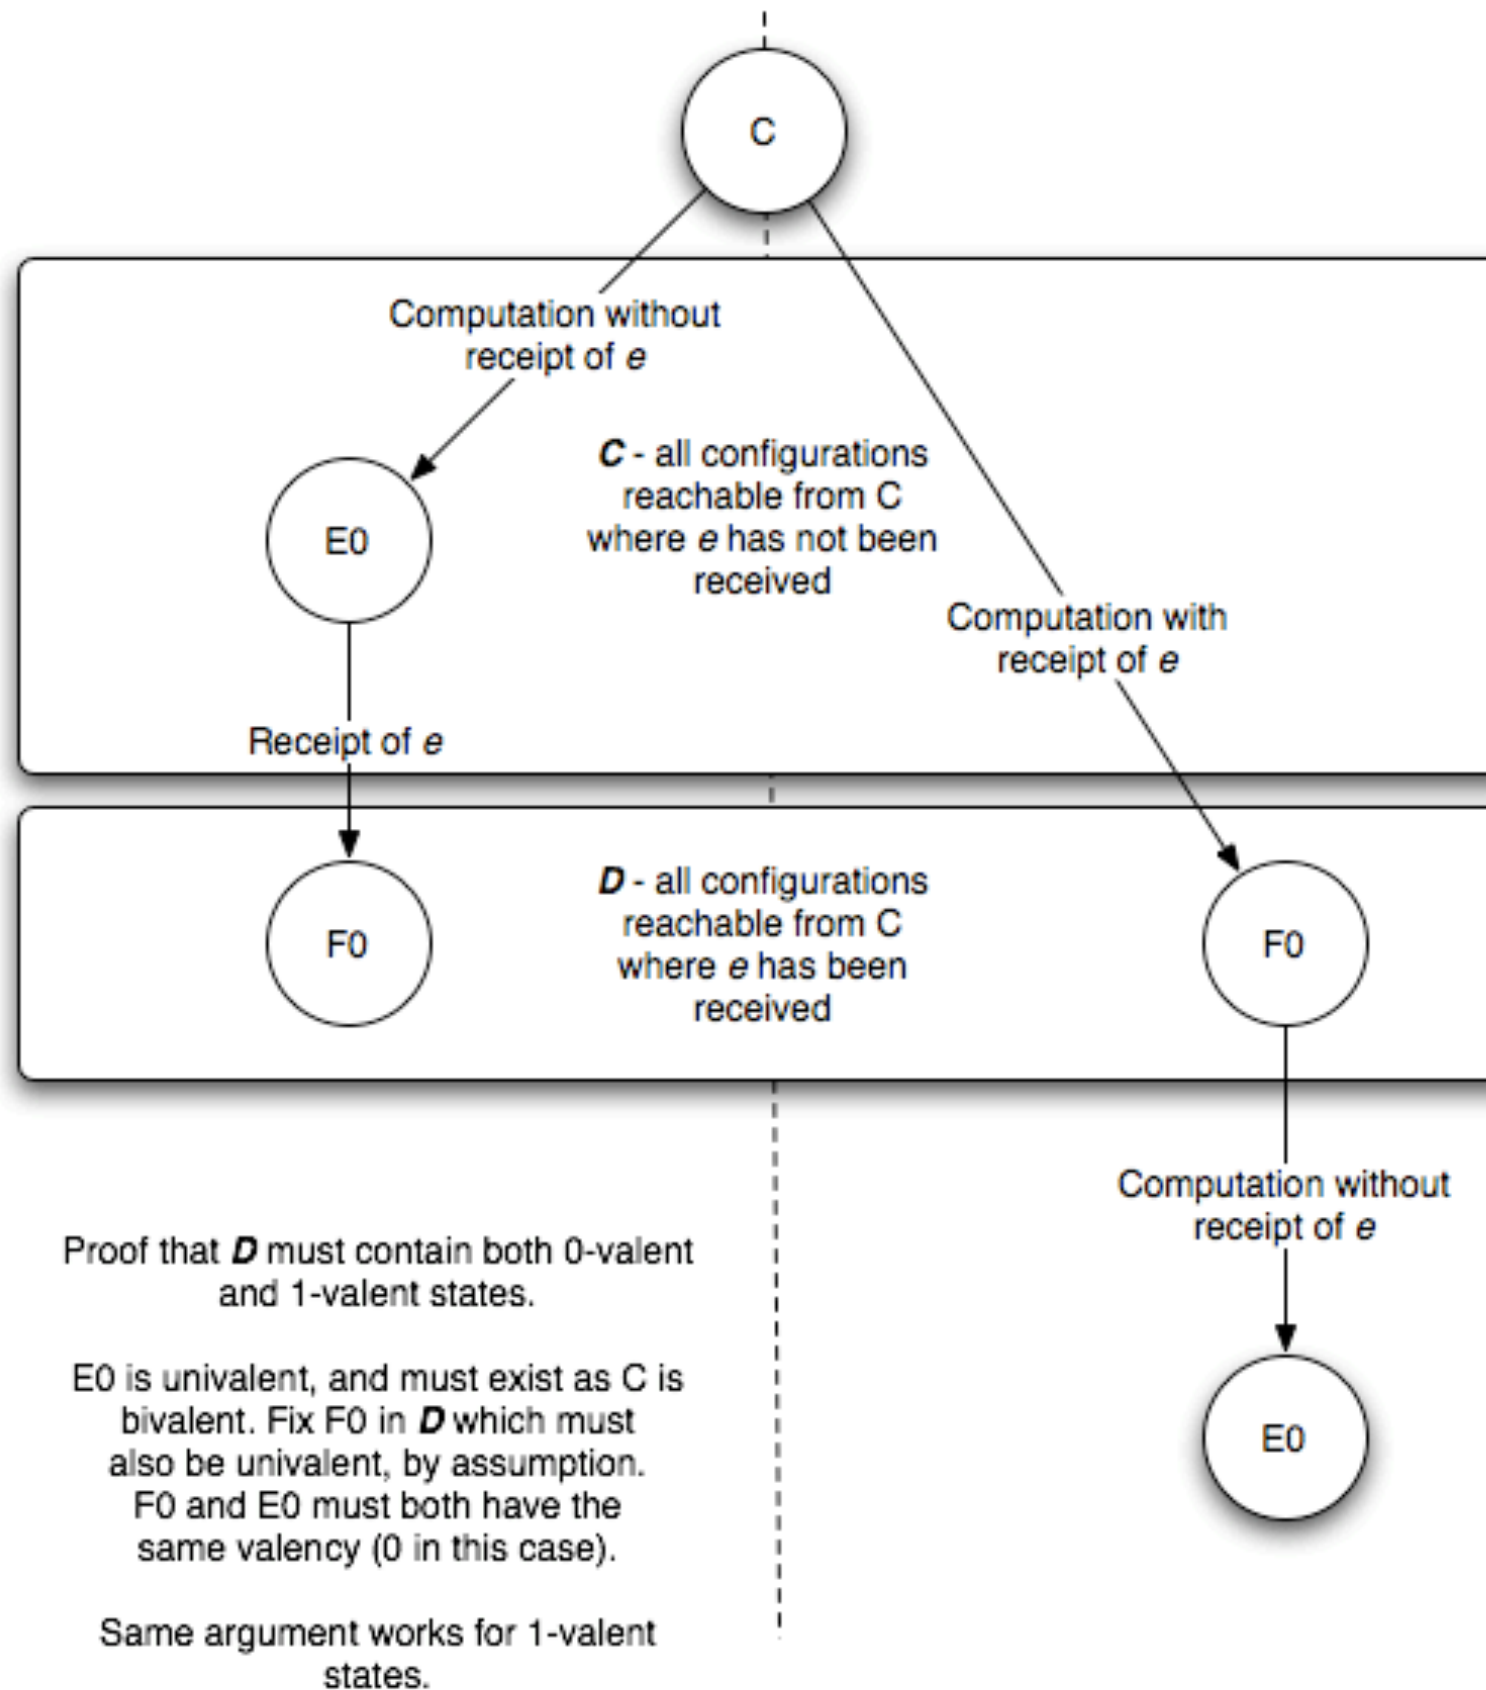
# Proving Step 2

- See <https://www.the-paper-trail.org/post/2008-08-13-a-brief-tour-of-flp-impossibility/>
- Good explanation of the proof
- Much more easier to understand than the paper



# Terms

- C: starting bivalent config
- R: configs where  $e$  has not been received
- D: configs where  $e$  has been received

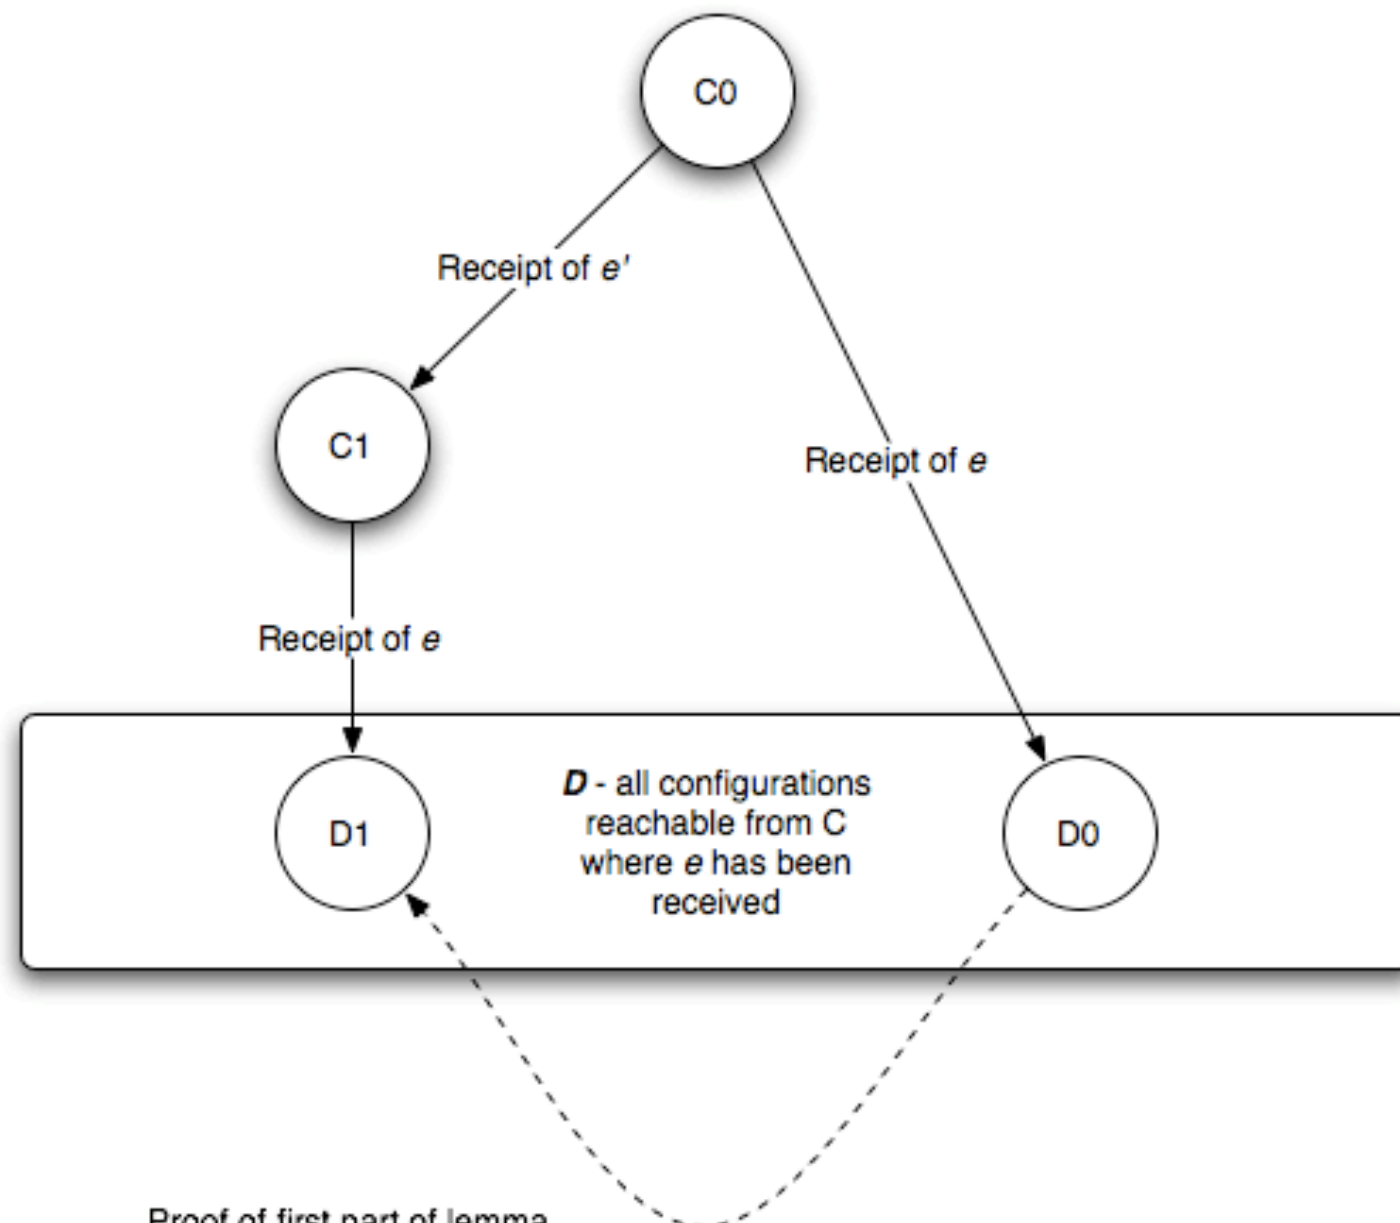


**Figure 2.1**



# Established so far..

- That  $D$  must have both 0-configs and 1-configs
- Assumption:  $D$  does not have bivalent configs



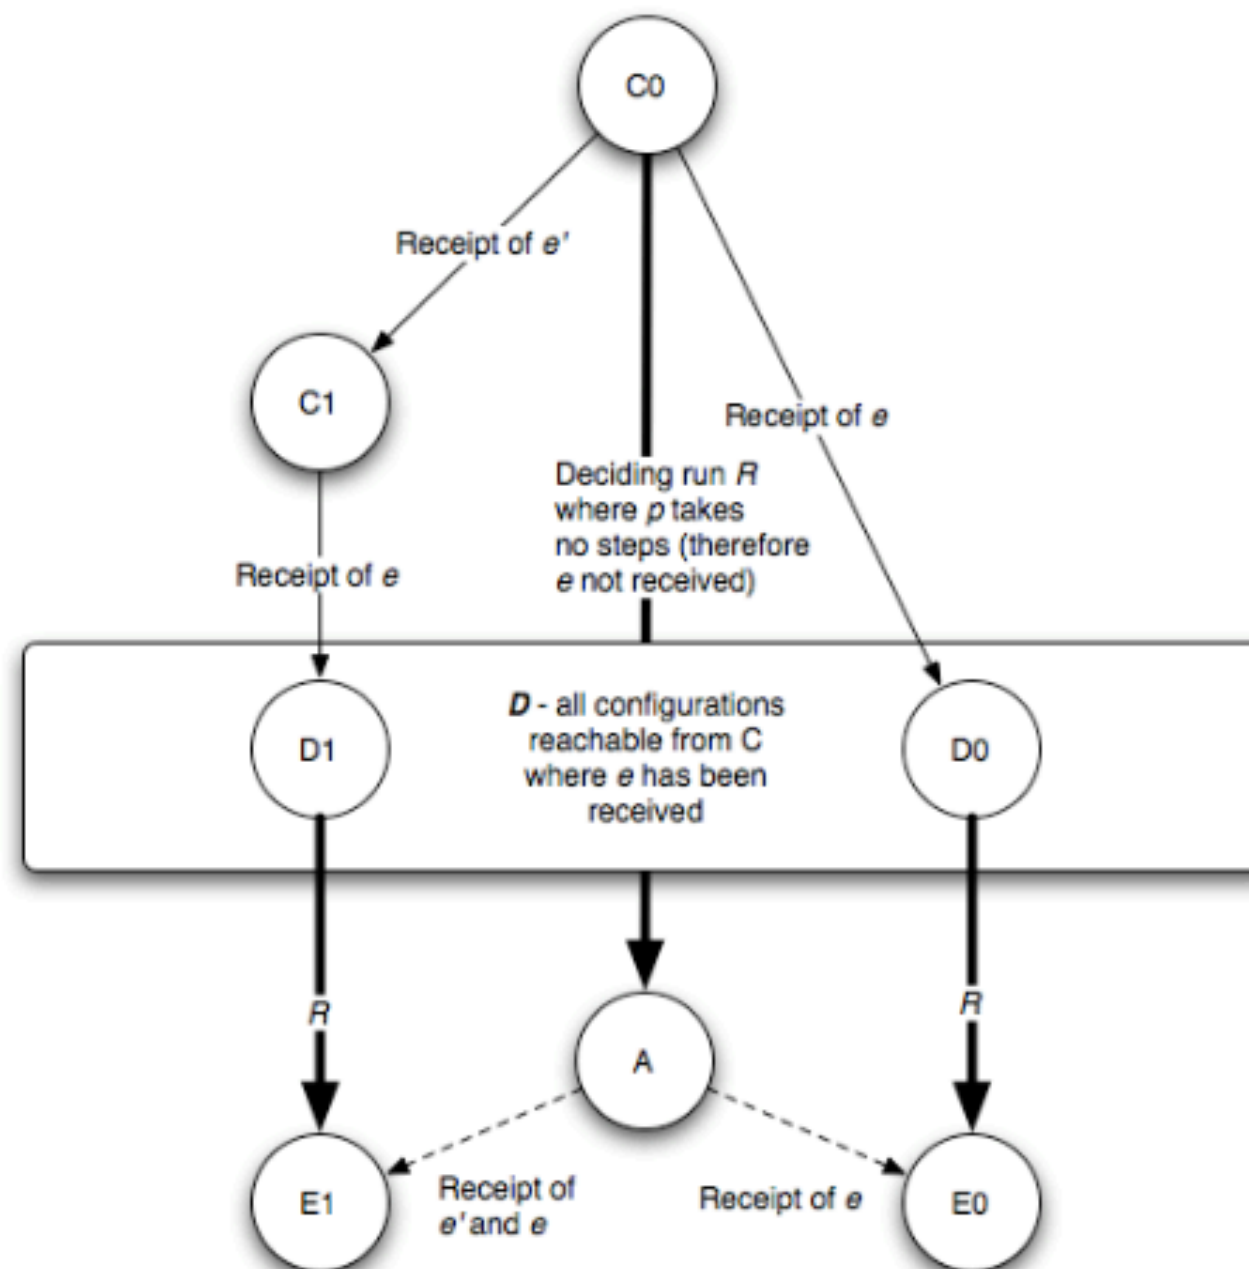
Proof of first part of lemma  
that **D** must contain bivalent state.

By assumption, D0 and D1 are  
have different univalencies. However  
if messages  $e$  and  $e'$  don't go to the  
same process they may be applied  
in any ordering and all processes  
must end in same state (as they have  
seen the same order of messages locally).

Receipt of  $e'$  if  
 $e$  and  $e'$  go to  
different  
processes

This is a contradiction as now D0 is bivalent.





Proof of second part of lemma, where  $e'$  and  $e$  are addressed to the same process  $p$ :  
 $E0$  and  $E1$  are 0- and 1-valent respectively.

$A$  is a univalent state reached by a deciding run from  $C0$  where the process  $p$  for which  $e'$  and  $e$  are intended takes no steps (as if it had failed).

But at  $A$  either  $e'$  then  $e$  or just  $e$  can be received by  $p$  which places it in one of two univalent states. But  $A$  is itself univalent, and so this is a contradiction.

# Constructing admissible runs which never decide

- We start from a bivalent configuration
- Let  $e$  be the message to be processed by process  $p$  next
- There is a bivalent configuration  $C'$  reachable from this config where  $e$  is the last message applied
- So we move from bivalent configs to other bivalent configs, and no decision is ever reached



# Paxos

- P1. An acceptor must accept the first proposal that it receives.
  - An acceptor can accept multiple values
- P2. If a proposal with value  $v$  is chosen, then every higher-numbered proposal that is chosen has value  $v$ .
- P2a. If a proposal with value  $v$  is chosen, then every higher-numbered proposal accepted by any acceptor has value  $v$ .
- P2b. If a proposal with value  $v$  is chosen, then every higher-numbered proposal issued by any proposer has value  $v$ .

# Paxos

- P2c. For any  $v$  and  $n$ , if a proposal with value  $v$  and number  $n$  is issued, then there is a set  $S$  consisting of a majority of acceptors such that either (a) no acceptor in  $S$  has accepted any proposal numbered less than  $n$ , or (b)  $v$  is the value of the highest-numbered proposal among all proposals numbered less than  $n$  accepted by the acceptors in  $S$ .
- P1a . An acceptor can accept a proposal numbered  $n$  iff it has not responded to a prepare request having a number greater than  $n$ .



# The Paxos Algorithm

- Phase 1. (a) A proposer selects a proposal number  $n$  and sends a prepare request with number  $n$  to a majority of acceptors.
- (b) If an acceptor receives a prepare request with number  $n$  greater than that of any prepare request to which it has already responded, then it responds to the request with a promise not to accept any more proposals numbered less than  $n$  and with the highest-numbered proposal (if any) that it has accepted.

# The Paxos Algorithm

- Phase 2. (a) If the proposer receives a response to its prepare requests (numbered  $n$ ) from a majority of acceptors, then it sends an accept request to each of those acceptors for a proposal numbered  $n$  with a value  $v$ , where  $v$  is the value of the highest-numbered proposal among the responses, or is any value if the responses reported no proposals.
- (b) If an acceptor receives an accept request for a proposal numbered  $n$ , it accepts the proposal unless it has already responded to a prepare request having a number greater than  $n$ .